

西南交通大学
Southwest Jiaotong University

计算机视觉及深度学习

卷积神经网络实验

--基于Tensorflow

教师：龚勋

Email：xgong@swjtu.edu.cn

卷积函数

`tf.nn.conv2d(input, filter, strides, padding,`

`use_cudnn_on_gpu=None, name=None)`

`tf.nn.depthwise_conv2d(input, filter, strides, padding,`

`name=None)`

`tf.nn.separable_conv2d(input, depthwise_filter, pointwise_filter,`

`strides, padding, name=None)`

tf.nn.conv2d

```
tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')
```

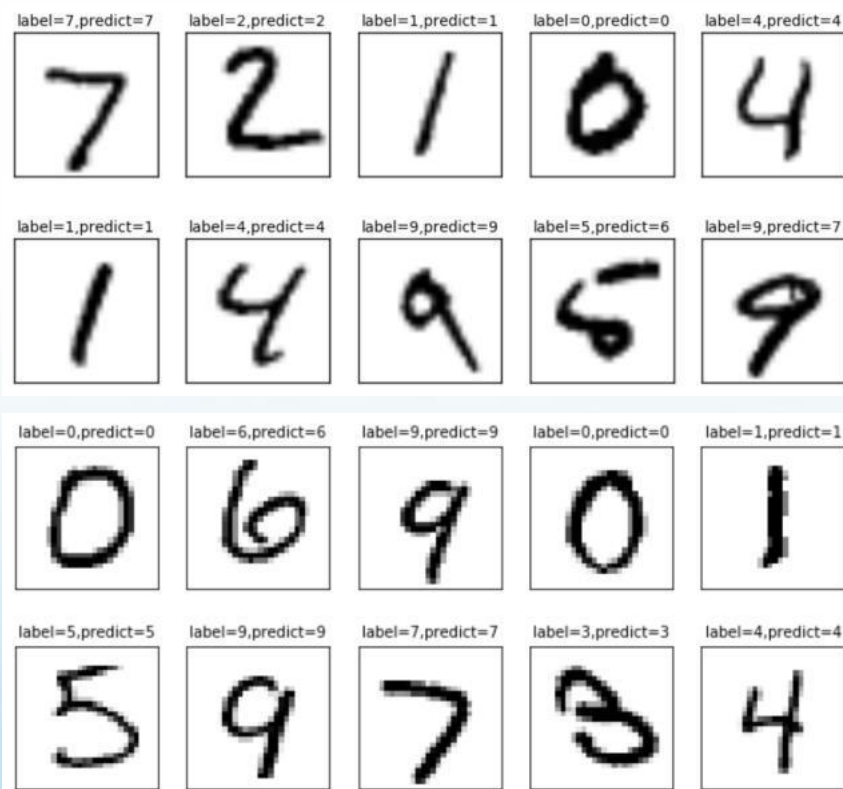

- **input**: 需要做卷积的输入数据。注意：这是一个4维的张量（[batch, in_height, in_width, in_channels]），要求类型为float32或float64其中之一。
- **filter**: 卷积核。[filter_height, filter_width, in_channels, out_channels]
- **strides**: 图像每一维的步长，是一个一维向量，长度为4
- **padding**: 定义元素边框与元素内容之间的空间。"SAME"或"VALID"，这个值决定了不同的卷积方式。当为"SAME"时，表示边缘填充，适用于全尺寸操作；当为"VALID"时，表示边缘不填充。
- **use_cudnn_on_gpu**: bool类型，是否使用cudnn加速
- **name**: 该操作的名称

池化函数

最大池化: `tf.nn.max_pool(value, ksize, strides, padding, name=None)`

平均池化: `tf.nn.avg_pool(value, ksize, strides, padding, name=None)`

基础实验1：MNIST手写数字识别



MNIST 数据集来自美国国家标准与技术研究所, National Institute of Standards and Technology (NIST).

数据集由来自 250 个不同人手写的数字构成, 其中 50% 是高中学生, 50% 来自人口普查局 (the Census Bureau) 的工作人员

训练集: 55000, 验证集: 5000, 测试集: 10000

<http://yann.lecun.com/exdb/mnist>www.swjtu.edu.cn

数据库读取





- MNIST 数据集可在网上获取 <http://yann.lecun.com/exdb/mnist/>
- TensorFlow提供了数据集读取方法

```
import numpy as np
import tensorflow as tf
import tensorflow.examples.tutorials.mnist.input_data as
input_data
mnist = input_data.read_data_sets('./data', one_hot=True)
```

is deprecated and will be removed in a future version.

MNIST手写数字识别数据集

MNIST数据集文件在读取时如果指定目录下不存在，则会自动去下载，需等待一定时间如果已经存在了，则直接读取

data > mnist			
Name	Date modified	Type	Size
 t10k-images-idx3-ubyte.gz	6/17/2017 5:05 PM	WinRAR 压缩文件	1,611 KB
 t10k-labels-idx1-ubyte.gz	6/17/2017 5:05 PM	WinRAR 压缩文件	5 KB
 train-images-idx3-ubyte.gz	6/17/2017 5:05 PM	WinRAR 压缩文件	9,681 KB
 train-labels-idx1-ubyte.gz	6/17/2017 5:05 PM	WinRAR 压缩文件	29 KB

MNIST手写数字识别数据集

```
In [3]: print(mnist.train.num_examples,mnist.validation.num_examples,
              mnist.test.num_examples)
print(mnist.train.images.shape,mnist.train.labels.shape)
print(mnist.validation.images.shape,mnist.validation.labels.shape)
print(mnist.test.images.shape,mnist.test.labels.shape)
```

```
55000 5000 10000
(55000, 784) (55000, 10)
(5000, 784) (5000, 10)
(10000, 784) (10000, 10)
```

训练、验证、测试集数据量


784=28*28, 数据维数

10: 标签维数


显示图像

```
%matplotlib inline
import matplotlib.pyplot as plt
def plot_image(image):
    plt.imshow(image.reshape(28,28), cmap='binary')
    plt.show()
```

```
plot_image(mnist.train.images[1])
```



```
plot_image(mnist.train.images[20000])
```




数据集划分

将单个数据集拆分为一个训练集和一个测试集



通过将数据集划分为三个子集，可以大幅降低过拟合的发生几率：



使用验证集评估训练集的效果。

在模型“通过”验证集之后，使用测试集再次检查评估结果

回到MNIST

- 验证集

```
print('image: ', mnist.validation.images.shape)  
print('labels: ', mnist.validation.labels.shape)
```

```
image: (5000, 784)
```

```
labels: (5000, 10)
```

- 测试集

```
print('image: ', mnist.test.images.shape)  
print('labels: ', mnist.test.labels.shape)
```

```
image: (10000, 784)
```

```
labels: (10000, 10)
```

采用线性模型构建优化方程

data

```
x = tf.placeholder(tf.float32, [None, 784], name="X")
```

```
y = tf.placeholder(tf.float32, [None, 10], name="Y")
```

parameters

```
W = tf.Variable(tf.random_normal([784, 10]), name="W")
```

```
b = tf.Variable(tf.zeros([10]), name="b")
```


模型


softmax分类

#model

`forward = tf.matmul(x,W) + b`

`pred = tf.nn.softmax(forward)`





```
sess = tf.Session()
init = tf.global_variables_initializer()
sess.run(init)
```

```
for epoch in range(train_epochs):
    for batch in range(total_batch):
        xs, ys = mnist.train.next_batch(batch_size)
        sess.run(optimizer, feed_dict={x:xs, y:ys})
    loss, acc = sess.run([loss_fun, accuracy],
                          feed_dict={x: mnist.validation.images,
                                      y: mnist.validation.labels})
    if epoch % display_step == 0:
        print("Epoch:%2d" % epoch, "loss:%.7f" % loss, "acc:%.7f" % acc)
print("Done! final acc = %f" % acc)
```

```
Epoch: 0 loss:5.3242607 acc:0.2898000
Epoch: 3 loss:2.1194503 acc:0.6140000
Epoch: 6 loss:1.5092000 acc:0.7084000
Epoch: 9 loss:1.2429380 acc:0.7502000
Epoch:12 loss:1.0902200 acc:0.7758000
Epoch:15 loss:0.9887187 acc:0.7976000
```

训练结果

.....

Epoch:129 loss:0.4691718 acc:0.8956000

Epoch:132 loss:0.4666167 acc:0.8966000

Epoch:135 loss:0.4632896 acc:0.8972000

Epoch:138 loss:0.4608540 acc:0.8984000

Epoch:141 loss:0.4579205 acc:0.8984000

Epoch:144 loss:0.4551692 acc:0.8986000

Epoch:147 loss:0.4529768 acc:0.8990000

Done! final acc = 0.898200

修改Mnist识别--卷积神经网络

```
# hidden layer
H1_NN = 6
H2_NN = 12
H3_NN = 24
FC1_NN = 200 # fully connected layer

x_image = tf.reshape(x, [-1, 28, 28, 1])
h_conv1 = conv2d_layer(x_image, [6, 6, 1, H1_NN], [H1_NN], tf.nn.relu)
h_pool1 = max_pool_2x2(h_conv1) #output is 14*14
h_conv2 = conv2d_layer(h_pool1, [5, 5, H1_NN, H2_NN], [H2_NN], tf.nn.relu)
h_pool2 = max_pool_2x2(h_conv2) #output is 7*7
h_conv3 = conv2d_layer(h_pool2, [4, 4, H2_NN, H3_NN], [H3_NN], tf.nn.relu)
h_pool3 = max_pool_2x2(h_conv3) #output is 4*4

h_pool3_flat = tf.reshape(h_pool3, [-1, 4 * 4 * H3_NN])
fc1 = fc_layer(h_pool3_flat, 4 * 4 * H3_NN, FC1_NN, tf.nn.relu)
```

```
def weight_variable(shape):
    initial = tf.truncated_normal(shape, stddev=0.1)
    return tf.Variable(initial)

def bias_variable(shape):
    initial = tf.constant(0.1, shape=shape)
    return tf.Variable(initial)

def conv2d(x, W):
    return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')

def max_pool_2x2(x):
    return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
                           strides=[1, 2, 2, 1], padding='SAME')

def conv2d_layer(inputs,
                  kernel_shape,
                  output_dim,
                  activation=None):
    W = weight_variable(kernel_shape)
    b = bias_variable(output_dim)
    h_conv = conv2d(inputs, W) + b
    if activation is None:
        outputs = h_conv
    else:
        outputs = activation(h_conv)
    return outputs
```

Epoch:26 loss:0.1022041 acc:0.9830000

Epoch:29 loss:0.1176529 acc:0.9816000

click to expand output; double click to hide output :0.9800000

Epoch:35 loss:0.1898626 acc:0.9756000

Epoch:38 loss:0.1107501 acc:0.9822000

Epoch:41 loss:0.1163032 acc:0.9820000

Epoch:44 loss:0.1304544 acc:0.9818000

Epoch:47 loss:0.1031366 acc:0.9822000



d:/ckpt_tf/mnist\mnistModel

Done! final acc = 0.980200, time=131.12s


```
test_acc = sess.run(accuracy, feed_dict={x: mnist
                                         y: mnist})
print("Test Acc=%.5f" % test_acc)
```

Test Acc=0.98090


Overfitting



Dropout




TRAINING
rate=0.5



EVALUATION
rate=0

```
pkeep = tf.placeholder(tf.float32)
```

```
fc1_dp = tf.nn.dropout(fc1, pkeep)
```



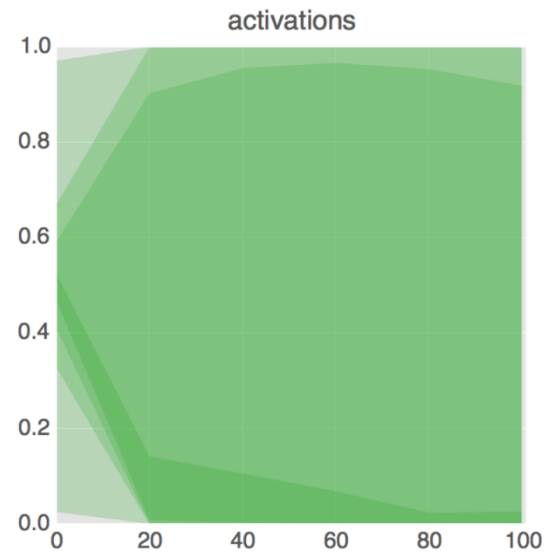

```
Epoch: 2 loss:0.0425114 acc:0.9870000 lr:0.0014147
Epoch: 5 loss:0.0361637 acc:0.9908000 lr:0.0006761
Epoch: 8 loss:0.0350688 acc:0.9920000 lr:0.0003525
Epoch:11 loss:0.0370960 acc:0.9920000 lr:0.0002106
Epoch:14 loss:0.0397607 acc:0.9922000 lr:0.0001485
+
Epoch:35 loss:0.0461600 acc:0.9922000 lr:0.0001002
Epoch:38 loss:0.0485143 acc:0.9918000 lr:0.0001001
Epoch:41 loss:0.0495552 acc:0.9922000 lr:0.0001000
Epoch:44 loss:0.0477311 acc:0.9918000 lr:0.0001000
Epoch:47 loss:0.0508552 acc:0.9916000 lr:0.0001000
d:/ckpt_tf/mnist/mnistModel
Done! final acc = 0.992000, time=137.62s
```

```
test_acc = sess.run(accuracy, feed_dict={x: mnist.test.images,
                                          y: mnist.test.labels,
                                          pkeep: 1.00})
print("Test Acc=%.5f" % test_acc)
```


```
Test Acc=0.99170
```

Batch Normal


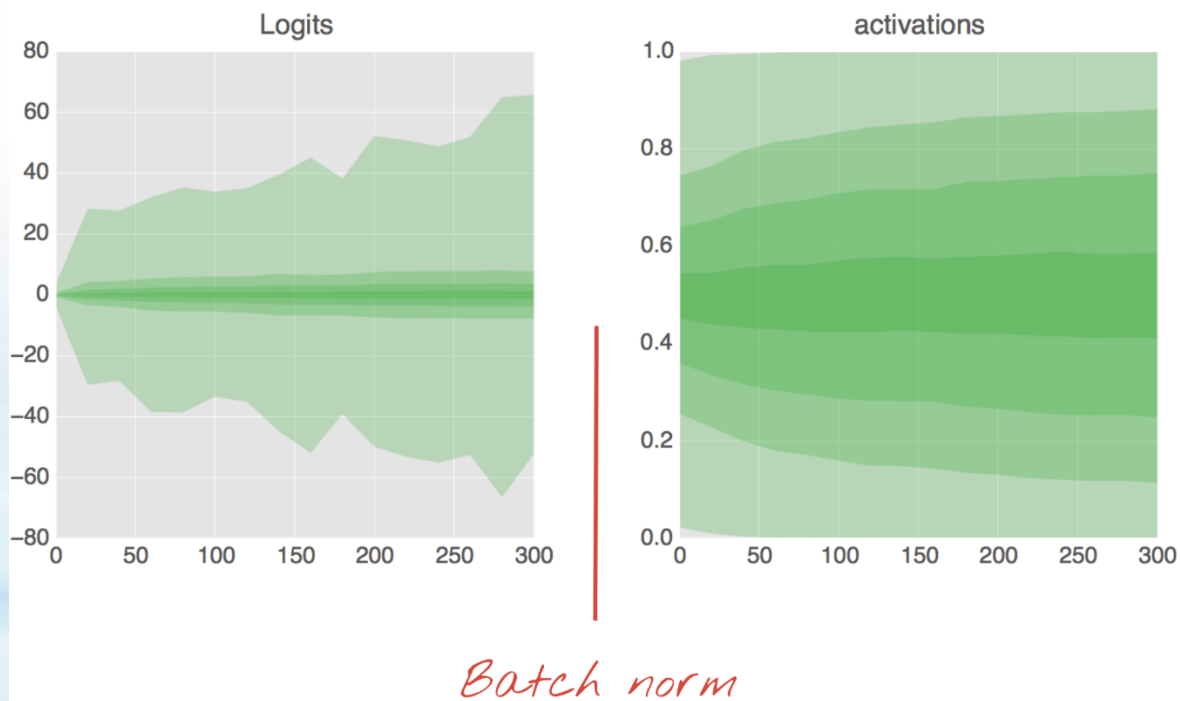
Without batch normalisation




boo-hoo



With batch normalisation (sigmoid)





```
def batchnorm(Ylogits, is_test, iteration, offset, convolutional=False):
    exp_moving_avg = tf.train.ExponentialMovingAverage(0.999, iteration) # adding the iter
    bnepsilon = 1e-5
    if convolutional:
        mean, variance = tf.nn.moments(Ylogits, [0, 1, 2])
    else:
        mean, variance = tf.nn.moments(Ylogits, [0])
    update_moving_averages = exp_moving_avg.apply([mean, variance])
    m = tf.cond(is_test, lambda: exp_moving_avg.average(mean), lambda: mean)
    v = tf.cond(is_test, lambda: exp_moving_avg.average(variance), lambda: variance)
    Ybn = tf.nn.batch_normalization(Ylogits, m, v, offset, None, bnepsilon)
    return Ybn, update_moving_averages
```

```
def conv2d_layer(inputs,
                  kernel_shape,
                  output_dim,
                  activation=None,
                  bn=True,
                  iteration=0,
                  if_test=False):
    W = weight_variable(kernel_shape)
    b = bias_variable(output_dim)
    B1 = bias_variable(output_dim)
    h_conv = conv2d(inputs, W) + b
    update_ema = 1.0
    if bn == True:
        h_convBN, update_ema = batchnorm(h_conv, if_test, iteration, B1, convolutional=True)
    else:
        h_convBN = h_conv

    if activation is None:
        outputs = h_convBN
    else:
        outputs = activation(h_convBN)
    return outputs, update_ema
```

基础实验2: Cifar10

The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

<http://www.cs.toronto.edu/~kriz/cifar.html>

airplane




automobile



bird



cat



deer



dog



frog



horse



ship











truck






Name

Size

 batches.meta	1 KB
 data_batch_1	30,309 KB
 data_batch_2	30,308 KB
 data_batch_3	30,309 KB
 data_batch_4	30,309 KB
 data_batch_5	30,309 KB
 readme.html	1 KB
 test_batch	30,309 KB



-  cifar-10-batches-py
-  cifar-10-binary.tar.gz
-  cifar-10-python.tar.gz

数据处理

定义全局变量

```
tf.app.flags.DEFINE_string('dbpath', './cifar-10-batches-  
py', '')
```

```
tf.app.flags.DEFINE_string('ckp_log_path', './cifar10',  
'')
```

```
FLAGS = tf.app.flags.FLAGS
```

```
def checkDb():
```

```
    if not os.path.exists(FLAGS.dbpath):
```

```
        print("DB does not exist!")
```

```
        exit(0)
```

```
    else:
```

```
        print(FLAGS.dbpath)
```

读取数据

```
def load_CIFAR_data(dbpath):  
    images_train = []  
    labels_train = []  
    for i in range(5):  
        fn = os.path.join(dbpath, 'data_batch_%d' % (i+1))  
        print(fn)  
        cur_img, cur_label = load_CIFAR_batch(fn)  
        images_train.append(cur_img)  
        labels_train.append(cur_label)  
    Xtrain = np.concatenate(images_train)  
    Ytrain = np.concatenate(labels_train)  
    Xtest, Ytest = load_CIFAR_batch(os.path.join(dbpath, 'test_batch'))  
    return Xtrain, Ytrain, Xtest, Ytest
```


读取一个文件

```
import pickle as p
```

```
def load_CIFAR_batch(filename):  
    with open(filename, 'rb') as f:  
        # image size 32*32*3  
        data_dict = p.load(f, encoding='bytes')  
        images = data_dict[b'data']  
        labels = data_dict[b'labels']  
        # original shape order: BCWH  
        images = images.reshape(10000, 3, 32, 32)  
        # covert to BWHC for TF  
        images = images.transpose(0, 2, 3, 1)  
        labels = np.array(labels)  
    return images, labels
```

主流程

```
if __name__ == '__main__':  
    # 1 load data  
    checkDb()  
    Xtrain, Ytrain, Xtest, Ytest = load_CIFAR_data(FLAGS.dbpath)  
    print(Xtrain.shape, Ytrain.shape)  
    print(Xtest.shape, Ytest.shape)  
    plot_images_labels(Xtrain, Ytrain, 0, 25, "images")  
    # 2 lable to OneHot  
    Ytrain_onehot = toOneHotCode(Ytrain)  
    Ytest_onehot = toOneHotCode(Ytest)  
    # 3 normalization  
    XtrainNorm = Xtrain.astype("float32") / 255.0  
    XtestNorm = Xtest.astype("float32") / 255.0  
    # 4 train  
    train(XtrainNorm, Ytrain_onehot, XtestNorm, Ytest_onehot)
```

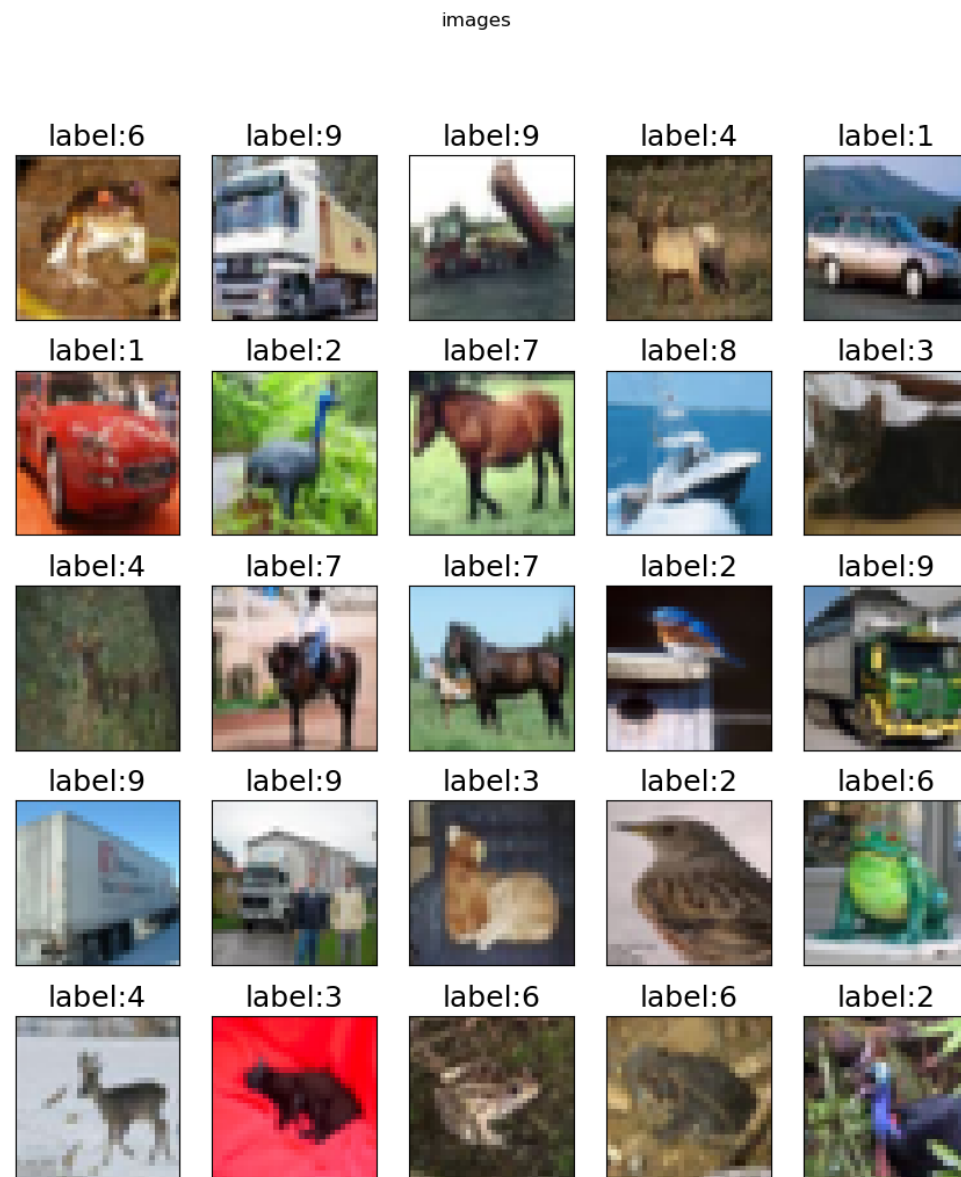
显示图像

```
def plot_images_labels(images, labels,
                       start, num=10,
                       caption=""):

    fig = plt.gcf()
    fig.set_size_inches(20,22)
    if num > 25:
        num = 25
    for i in range(0, num):
        ax = plt.subplot(5,5,i+1)
        ax.imshow(images[start], cmap='bgr')
        title = 'label:' + str(labels[start])
        ax.set_title(title, fontsize=18)
        ax.set_xticks([])
        ax.set_yticks([])
        start += 1

    plt.suptitle(caption)

plt.show()
```



数据预处理

One-Hot-Code

```
def toOneHotCode(ori):  
    encoder = OneHotEncoder(sparse=False)  
    sample = [[0],[1],[2],[3],[4],[5],[6],[7],[8],[9]]  
    encoder.fit(sample)  
    ori_re = ori.reshape(-1,1)  
    one = encoder.transform(ori_re)  
    return one
```

Normalization

```
XtrainNorm = Xtrain.astype("float32") / 255.0  
XtestNorm = Xtest.astype("float32") / 255.0
```



```
def net(input, pkeep):
```

```
#####model#####
```

```
# hidden layer
```

```
H1_NN = 16;H2_NN = 32;H3_NN = 64;FC1_NN = 128
```

```
h_conv1 = conv2d_layer(input, [3, 3, 3, H1_NN], [H1_NN], tf.nn.relu)
```

```
h_pool1 = max_pool_2x2(h_conv1) # output is 14*14
```

```
h_conv2 = conv2d_layer(h_pool1, [3, 3, H1_NN, H2_NN], [H2_NN], tf.nn.relu)
```

```
h_pool2 = max_pool_2x2(h_conv2) # output is 7*7
```

```
h_conv3 = conv2d_layer(h_pool2, [3, 3, H2_NN, H3_NN], [H3_NN], tf.nn.relu)
```

```
h_pool3 = max_pool_2x2(h_conv3) # output is 4*4
```

```
h_pool3_flat = tf.reshape(h_pool3, [-1, 4 * 4 * H3_NN])
```

```
fc1 = fcn_layer(h_pool3_flat, 4 * 4 * H3_NN, FC1_NN, tf.nn.relu)
```

```
fc1_dp = tf.nn.dropout(fc1, pkeep) ###
```

```
forward = fcn_layer(fc1_dp, FC1_NN, 10, None)
```

```
pred = tf.nn.softmax(forward)
```

```
return forward, pred
```

训练

```
for epoch in range(train_epochs):
    # shuffle at the beginning of each epoch
    Xtrain, Ytrain = shuffle(Xtrain, Ytrain)
    for batch in range(total_batch):
        xs, ys, idx_in_epoch = next_batch(Xtrain, Ytrain, batch_size, idx_in_epoch)
        _, loss, lr = sess.run([optimizer, loss_fun, learning_rate],
                                feed_dict={x: xs, y: ys, pkeep: 0.75, step: iteration})


        iteration += 1
    # validation
    acc = sess.run(accuracy, feed_dict={x: Xtest,
                                         y: Ytest,
                                         pkeep: 1.0, step: iteration})

    if (epoch + 1) % display_step == 0:
        print("Epoch:%2d" % epoch, "loss:{:.7f}".format(loss),
              "lr:{:.7f}" % lr, "acc:{:.7f}".format(acc))
    if (epoch + 1) % save_step == 0:
        fn = os.path.join(FLAGS.ckp_log_path, 'cifar10Model')
        print(fn)
        saver.save(sess, fn, global_step=(epoch + 1))
print("Done! final acc = %f, time=%.2fs" % (acc, time() - start))
```


next_batch

参考mnist.py

```
def next_batch(X, Y, batch_size, _index_in_epoch=0, shuffle=True):
    start = _index_in_epoch;    _num_examples = X.shape[0]
    if start + batch_size > _num_examples:
        ...
    else:
        _index_in_epoch += batch_size
        end = _index_in_epoch
        return X[start:end], Y[start:end], _index_in_epoch
```



Epoch: 9	loss:0.5913961	lr:0.0007292	acc:0.7131000
Epoch:10	loss:0.7697335	lr:0.0006382	acc:0.7150000
Epoch:11	loss:0.6140158	lr:0.0005604	acc:0.7122000
Epoch:12	loss:0.3961719	lr:0.0004938	acc:0.7173000
Epoch:13	loss:0.7391727	lr:0.0004369	acc:0.7189000
Epoch:14	loss:0.3909289	lr:0.0003881	acc:0.7154000
Epoch:15	loss:0.4020708	lr:0.0003465	acc:0.7185000
Epoch:16	loss:0.3959954	lr:0.0003108	acc:0.7217000
Epoch:17	loss:0.4575052	lr:0.0002803	acc:0.7191000
Epoch:18	loss:0.2468506	lr:0.0002543	acc:0.7202000
Epoch:19	loss:0.6702175	lr:0.0002320	acc:0.7199000
Epoch:20	loss:0.1293726	lr:0.0002129	acc:0.7217000
Epoch:21	loss:0.3751268	lr:0.0001966	acc:0.7232000

Go Deep

```
def net(x, y, pkeep):
    #####model#####
    # hidden layer
    H1_NN = 16
    H2_NN = 32
    H3_NN = 64
    H4_NN = 128
    FC1_NN = 256 # fully connected layer


    h_conv1 = conv2d_layer(x, [3, 3, 3, H1_NN], [H1_NN], tf.nn.relu)
    h_pool1 = max_pool_2x2(h_conv1) # output is 14*14
    h_conv2 = conv2d_layer(h_pool1, [3, 3, H1_NN, H2_NN], [H2_NN], tf.nn.relu)
    h_pool2 = max_pool_2x2(h_conv2) # output is 7*7
    h_conv3 = conv2d_layer(h_pool2, [3, 3, H2_NN, H3_NN], [H3_NN], tf.nn.relu)
    h_pool3 = max_pool_2x2(h_conv3) # output is 4*4
    print('h_pool3 shape: ' + str(h_pool3.shape))
    h_conv4 = conv2d_layer(h_pool3, [3, 3, H3_NN, H4_NN], [H4_NN], tf.nn.relu)

    h_pool4_flat = tf.reshape(h_conv4, [-1, 4 * 4 * H4_NN])
    fc1 = fcn_layer(h_pool4_flat, 4 * 4 * H4_NN, FC1_NN, tf.nn.relu)
    fc1_dp = tf.nn.dropout(fc1, pkeep) ###

    forward = fcn_layer(fc1_dp, FC1_NN, 10, None)
    pred = tf.nn.softmax(forward)
    return forward, pred
```

练习

- 将Cifar10数据读取及操作（next_batch）封装成一个类进行调用
- 增加网络深度，采用ResNet50进行网络修改，将准确率提升（>85%）



西南交通大学
Southwest Jiaotong University

Thank You !

